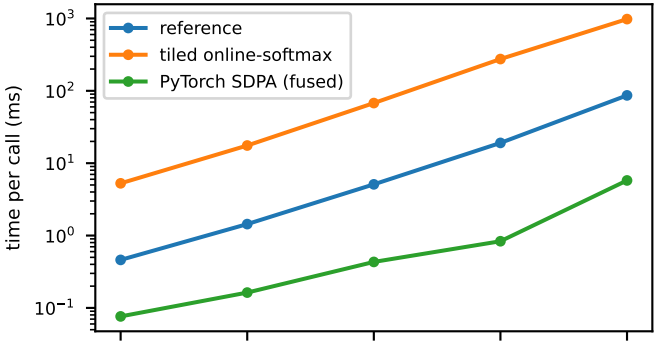
reference
tiled online-softmax
PyTorch SDPA (fused)
time per call (ms)
10^{3}
10^{2}
10^{1}
10^{0}
10^{-1}
2^{8}
2^{9}
2^{10}
2^{11}
2^{12}
sequence length T

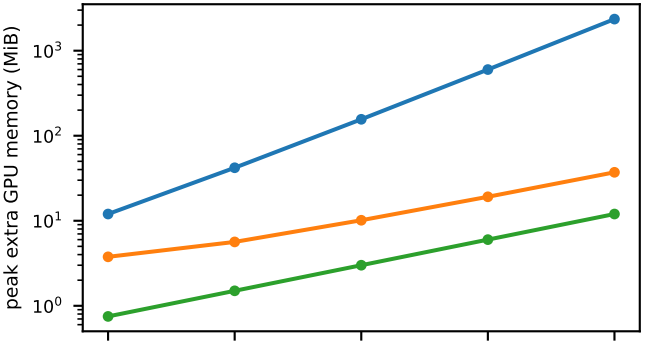
peak extra GPU memory (MiB)
10^{3}
10^{2}
10^{1}
10^{0}
2^{8}
2^{9}
2^{10}
2^{11}
2^{12}
sequence length T